

Operations ... SQL Specifics



pm jat @ daiict



Things that we talk about

- Result of SQL SELECT [SET or MULTI-SET] ?
- Nulls as special values in SQL
- Sub-queries in SQL
- EXISTS and NOT EXISTS in SQL
- Functions and Operators in SQL
- ORDER BY
- LIMIT and OFFSET



Operations ... SQL Specifics

Result of SQL SELECT [SET or MULTI-SET] ?



Result of SQL SELECT [SET or MULTI-SET] ?

- By Definition, relations are set; but implementations may permit duplicate tuples and such relations are called *bag (multi-set)*
- Normally stored relations (base) relations should still be sets, because most relations have Primary Key
- However SQL SELECT does not produce SET, their results are often bags
 - possibly because duplicate removal is expensive.
 - We can use DISTINCT keyword for returning sets
- On the other hand, SQL SET operation (UNION/INTERSECT/EXCEPT) produce a valid “SET”.
 - possibly because duplicate removal comes as natural outcome of executing set operations on relations.
 - But again, SQL allows producing bags too here through UNION ALL, EXCEPT ALL, and so



UNION/INTERSECT/EXCEPT ALL in SQL

- Compare result of following queries:

```
SELECT super_eno FROM employee; --Q1
```

```
SELECT mgr_eno FROM department; --Q2
```

```
SELECT super_eno FROM employee
```

```
UNION
```

```
SELECT mgr_eno FROM department; --Q3
```

```
SELECT super_eno FROM employee
```

```
UNION ALL
```

```
SELECT mgr_eno FROM department; --Q4
```



UNION ALL in SQL

- It just combines tuples from operand relations. The figure should depict the operation.

R	
super_eno	
105	
105	
null	
106	
106	
102	
106	
102	
101	

S	
mgr_eno	
102	
106	
104	
106	

R UNION S		R UNION ALL S	
super_eno		super_eno	
null		105	
106		105	
102		null	
104		106	
105		106	
		102	
		106	
		102	
		101	
		102	
		106	
		104	
		106	



INTERSECT ALL in SQL

- Common occurrences of tuples are included in the result.
- Let us say a tuple t appears m times in R and n times in S ; then $\min(m,n)$ times the tuple t would appear in the result.

R
super_eno
105
105
null
106
106
102
106
102
101

S
mgr_eno
102
106
104
106

R INTERSECT S		R INTERSECT ALL S
super_eno		super_eno
102		102
106		106
		106



EXCEPT ALL in SQL

- Occurrences of tuples are subtracted from the right operand relation.
- Let us say a tuple t appears m times in R and n times in S ; then $m-n$ times, the tuple t would appear in the result.

R EXCEPT S		R EXCEPT ALL S
super_eno		super_eno
null		null
101		101
105		105
		105
		102
		106

R
super_eno
105
105
null
106
106
102
106
102
101

S
mgr_eno
102
106
104
106



Operations ... SQL Specifics

Nulls as special values in SQL



Meaning of Null Values for an attribute

- NULLs require special treatment in different usage contexts in SQL.
- Because an attribute having NULL could mean either of the following-
 - Value is unknown, that is not available right now
 - No value for the attribute applicable to the tuple
- We have already seen how it is treated in aggregate operations.
 - Hope that treatment makes sense to you in this context of usage!
- Let us see, how it is interpreted when it appears in
 - Arithmetic expressions
 - Logical expressions in predicate



Meaning of Null Values for an attribute

- For example, how following expressions would be evaluated for tuples where salary is NULL?
- Arithmetic expression
`SELECT e.salary*1.1 from employee as e;`
- Logical expression in predicate
`Select * from employee as e where e.salary > 50000;`
`Select * from employee as e where e.salary > 50000 and dno=4;`
- Let us take one by one



“Null Value” in arithmetic expressions

- Arithmetic expressions (+,-,*,/) involving null values result null
- There for expression “e.salary*1.1” would evaluate to NULL for employees having NULL for salary!
- Make sense?



Interpretation of “Null Value” in predicate

- Now suppose we have following query
Select * from employee as e where `e.salary > 50000`;
- How these expressions would be evaluated for tuples where salary is NULL?
- Actually Interpret `e` here as tuple variable that ranges over all tuples of employee relations.
- For example, consider a query: `select * from r where <condition>`
where r can be any “relational expression”; i.e. `r1 join r2` or so
- Let us say a query is executed as following:
for tuple `t ∈ r`
if (`cond` is `true`)
add `t` to resultset



Interpretation of “Null Value” in predicate

- In predicate expressions, the occurrence of NULL is read as “UNKNOWN”
- That means, in expressions like $t.a < 10$, $t.a$ is taken as “UNKNOWN” when it happens to be NULL.
- For this purpose, SQL defines “UNKNOWN” as the third truth value (in addition to TRUE and FALSE) for evaluating logical expressions; and
- While evaluating WHERE clause tuples with UNKNOWN or FALSE truth values are not included in the resultset
- So, now we have the answer to the question that how expressions in the following query would be evaluated for tuples where salary is NULL.
Select * from employee as e where $e.salary > 50000$;



Truth values for UNKNOWN

- NOT
 - NOT UNKNOWN → UNKNOWN
- AND
 - TRUE AND UNKNOWN → UNKNOWN
 - FALSE AND UNKNOWN → FALSE
 - UNKNOWN AND UNKNOWN → UNKNOWN
- OR
 - TRUE OR UNKNOWN → TRUE
 - FALSE OR UNKNOWN → UNKNOWN
 - UNKNOWN OR UNKNOWN → UNKNOWN



Null Values and Comparisons

- Following query will not include any tuple where either of value in NULL irrespective value in other attribute

```
SELECT * FROM EMPLOYEE WHERE  
bdate < DATE '2001-01-01' AND salary > 30000
```

- Following query will not include a tuple only when both are NULL, if one of attribute meets the condition then it will get included in result

```
SELECT * FROM EMPLOYEE WHERE  
bdate < DATE '2001-01-01' OR salary > 30000
```




Null Values and Comparisons – IS NULL

- Following will not give desired result. Why? -

```
SELECT * FROM employee  
      WHERE super_eno = NULL;
```

- This is so because Null = Null is also UNKNOWN. For checking an attribute for having a NULL value, SQL provides IS NULL (and IS NOT NULL)
- We write as following for such situations –

```
SELECT * FROM employee  
      WHERE super_eno IS NULL;
```



Subquery in SQL

- A Query that is part of another query is *subquery*. A subquery may also have subquery, and so forth upto any level

SELECT _____ (**SELECT** _____)

- A subquery in SQL is written as a *query expression* enclosed in parentheses, and is in following form-
" (**SELECT** . . . **FROM** ...) "
as a part of some existing query
- **Result of sub-query is again a relation;**



Subquery in FROM clause

- Here is an example

```
select ename, dno, salary from employee natural join  
  (select eno from works_on natural join project  
   where pno=1) as r;
```

```
select r1.*, m.salary as manager_salary from  
  (select e.eno as enum, e.ename as emp_name,  
         e.salary as emp_salary,  
         d.mgr_eno as manager  
   FROM employee e join department d  
   on e.dno=d.dno) as r1  
  join employee m on (manager=eno)  
 where emp_salary > m.salary;
```



Subquery in FROM clause

```
select eno, ename, salary from employee  
natural join  
(select eno from employee  
except  
select mgr_eno from department) as r;
```



Subquery in WHERE clause

```
select * from employee where eno not in  
      (select eno from works_on);
```

- Sub query in FROM clause is never a problem. However, when subquery in where clause, there needs to be few cautions to be noted!
- Recall that where clause is evaluated for every tuple of operand relation?
 - Does it mean; subquery here needs to be executed N times?
 - Note subquery here may also be called as inner query!



Execution of Subquery

- **SUB-Query may not execute for every tuple of outer query**

- Consider another query-

```
SELECT * FROM student WHERE  
progid IN (SELECT pid FROM program  
          WHERE did = 'EE' );
```

- Typically, after execution of inner query, outer query may be translated to:

```
SELECT * FROM student WHERE  
progid IN (BEC, BEE);
```

- However this optimization may not be possible when you have
“*correlated sub-query*”



Correlated Sub-Queries

- When inner query refers to a tuple of the outer query then it is a correlated sub-query.
- Consider the following query –

```
SELECT eno, ename FROM employee as e WHERE salary >  
(SELECT AVG(salary) FROM employee WHERE dno = e.dno)
```
- Are you able to determine, what does it compute?
“List employees, whose salary is more than the department average”



Execution of Correlated Sub-Queries

- Consider same query

```
SELECT eno, ename FROM employee as e  
WHERE salary > (SELECT AVG(salary)  
FROM employee WHERE dno = e.dno)
```

- Logically, it is as the following:
For each tuple of the outer query, execute the inner query.
- Note that it can not be executed once for all tuples of the outer query, as the case be with the un-related inner query, and we have to execute **SUB-Query for every tuple of the outer query**
- This is an identified problem with correlated sub-queries.



Correlated Sub-Queries could be expensive to execute – therefore should be avoided

- Correlated queries are expensive to execute, and can be avoided; for example the previous example

```
SELECT eno, ename FROM employee as e  
WHERE salary > (SELECT AVG(salary)  
FROM employee WHERE dno = e.dno)
```

- can be re-written as-

```
SELECT eno, ename, salary FROM employee as e NATURAL JOIN (SELECT  
dno, AVG(salary) as avg_sal FROM employee GROUP BY dno) as av  
WHERE salary > av.avg_sal;
```



more Correlated Sub-queries

- List down employees having salary greater than their immediate supervisors.
`select * from employee as e1 where e1.salary > (select salary from employee as e2 where e2.eno = e1.super_eno);`
- Select employees having dependents older than 18 years:
`SELECT * FROM employee AS e WHERE eno IN (SELECT eno FROM dependent AS d WHERE d.eno = e.eno AND age(d.bdate) > interval '18 years');`
- Attempt re-writting them without correlated query.



Compare a values with a bag of values (SQL)

- For example consider following two queries
[Find out employee who have salary greater some or all employees of dno = 4]

SELECT enon, ename FROM employee WHERE salary

> **SOME** (SELECT salary FROM employee WHERE dno = 4);

SELECT eno, ename FROM employee WHERE salary

> **ALL** (SELECT salary FROM employee WHERE dno = 4);



Compare a values with a bag of values (SQL)

- Note the equivalences:

SELECT enon, ename FROM employee WHERE salary

> **SOME** (SELECT salary FROM employee WHERE dno = 4); and

SELECT enon, ename FROM employee WHERE salary

> (SELECT **min**(salary) FROM employee WHERE dno = 4);

SELECT enon, ename FROM employee WHERE salary

> **ALL** (SELECT salary FROM employee WHERE dno = 4); and

SELECT enon, ename FROM employee WHERE salary

> (SELECT **max**(salary) FROM employee WHERE dno = 4);



Compare a values with a bag of values (SQL)

- Comparative operators could be, one of following-
 >SOME, >=SOME <=SOME, <SOME, =SOME, <>SOME
 >ALL, >=ALL, <=ALL, <ALL, =ALL, <>ALL



Sub-queries in Update statements

- **UPDATE** employee
 SET salary = salary * 1.1
 WHERE eno IN (...);
- **DELETE** employee
 WHERE eno = (...);



EXISTS and NOT EXISTS in SQL

- Checks for emptiness of a relation and returns true or false.
- **EXISTS (r)** can be interpreted as “is there some tuple exists in relation r”
- **EXISTS (r)** returning true says that argument relation r is not empty
- Similarly, **NOT EXISTS (r)** returning true says that argument relation r empty



Example EXISTS

- List employees who have dependents older than 18 years
- `SELECT * FROM employee AS e WHERE EXISTS (SELECT * FROM dependent AS d WHERE d.eno = e.eno AND age(d.dob) > interval '18 years');`



SQL- EXISTS and IN

- While they might appear to be serving similar purposes, semantically are different.
- Both appear as part of predicate in WHERE clause of SELECT
- IN:
 - Syntax: **x IN (r)**
 - Meaning: checks existence of tuple x in relation r , if found returns true, other wise false. Normally x is a scalar value and r is a single column relation.
- EXISTS:
 - Syntax: **EXISTS (r)**
 - Meaning: checks if r is a non empty relation. Returns true if the relation has at least one tuple, otherwise false.
- In both above cases r is a *relational expression* resulting a relation.



Functions and Operators in SQL



Functions and Operators

- SQL provides various functions and operators that can be used to create a new attribute in resultant relations
- There are typically, type conversion, arithmetic operators, mathematical, and string manipulation operators and functions. For example: `substring`, `upper`, `lower`, `sqrt`, `ln`, etc.
- Details for PostgreSQL functions can be seen at:
http://intranet.daiict.ac.in/~pm_jat/postgres/html/functions.html.



Examples

```
SELECT eno,  
       fname || ' ' || minit || '. ' || lname AS name,  
       current_date - dob AS age FROM employee;
```

```
SELECT eno, hours*50 AS amount FROM works_on;
```

```
SELECT upper(ename) AS name, ln(salary) AS x FROM employee;
```

```
SELECT * FROM employee  
       WHERE upper(ename) = 'AMIT';
```

```
SELECT eno FROM dependent WHERE age(d.dob) > interval '18 years');
```



BETWEEN and LIKE in SQL

- **BETWEEN, LIKE** are used in predicate:

```
SELECT ... WHERE A BETWEEN 10 TO 20;
```

```
SELECT ... WHERE A1 LIKE '%IX%' OR A2 LIKE 'ABC%' OR A3 LIKE '%XYZ';
```

```
SELECT ... WHERE A1 LIKE '_X_%';
```

- Also: **NOT BETWEEN** and **NOT LIKE**.



Regular Expression Matching in PostgreSQL

- PostgreSQL also allows regular expression matching in string match using **IS SIMILAR TO <reg-ex>**



ORDER BY

- ORDER BY CLAUSE is used for ordering the resultant tuples of a SQL query.
- Following statements returns row-set from employee table, and rows are sorted based on salary. To order in descending order, we add DESC keyword after attribute name.

```
select * from employee order by salary;
```

```
select * from employee order by salary desc;
```

- Following statement returns row-set from employee table, and rows are sorted in ascending order of dno, and within dno all rows are sorted on salary in descending order-

```
select * from employee order by dno, salary desc;
```



LIMIT and OFFSET

- Examples below should be self explanatory
- Gives top three earners

```
select * from employee order by salary desc limit 3
```

- Gives next two earners after top 3

```
select * from employee order by salary desc  
offset 3 limit 2
```